
Ed Zitron on big tech, backlash, boom and bust: 'AI has taught us that people are excited to replace human beings'

Published by

Steve Rose

Estimated reading time

13 min

Added on

2026-01-19

Address

https://www.theguardian.com/technology/2026/jan/19/ed-zitron-on-big-tech-backlash-boom-and-bust-ai-has-taught-us-that-people-are-excited-to-replace-human-beings?CMP=GTUK_email

Ed Zitron on big tech, backlash, boom and bust: 'AI has taught us that people are excited to replace human beings'

If some time in an entirely possible future they come to make a movie about "how the AI bubble burst", Ed Zitron will doubtless be a main character. He's the perfect outsider figure: the eccentric loner who saw all this coming and screamed from the sidelines that the sky was falling, but nobody would listen. Just as Christian Bale portrayed [Michael Burry](#), the investor who predicted the 2008 financial crash, in [The Big Short](#), you can well imagine Robert Pattinson fighting Paul Mescal, say, to portray Zitron, the animated, colourfully obnoxious but doggedly detail-oriented Brit, who's become one of big tech's noisiest critics.

This is not to say the AI bubble *will* burst, necessarily, but against a tidal wave of AI boosterism, Zitron's blunt, brash scepticism has made him something of a cult figure. His tech newsletter, [Where's Your Ed At](#), now has more than 80,000 subscribers; his weekly podcast, [Better Offline](#), is well within the Top 20 on the tech charts; he's a regular dissenting voice in the media; and his [subreddit](#) has become a safe space for AI sceptics, including those within the tech industry itself – one user describes him as "a lighthouse in a storm of insane hypercapitalist bullshit".

Zitron first started looking into generative AI in 2023, a year after the industry-shaking launch of OpenAI's ChatGPT. "The more I looked, the more confused I became, because on top of the fact that large language models (LLMs) very clearly did not do the things that people were excited about, they didn't have any path to doing them either," he says. "Nothing I found made any suggestion that this was a real business at all, let alone something that would supposedly change the world."

He's talking over videocall from his office in Las Vegas, dressed in a red hoodie, surrounded by framed pop-culture prints and American sports memorabilia. And boy can Zitron talk. As listeners to Better Offline will know, the 39-year-old is a prodigious speaker – adept at extended monologues, putting his point of view across in accessible, often cheeky language, peppered with facts, statistics, analogies and a fair few expletives, in a London accent that only accentuates his position as a Silicon Valley contrarian – someone who drops his Ts when he says "datacentres".

Explaining Zitron's thesis about why generative AI is doomed to fail is not simple: last year he wrote [a 19,000-word essay](#), laying it out. But you could break it down into two, interrelated parts. One is the actual efficacy of the technology; the other is the financial architecture of the AI boom. In Zitron's view, the foundations are shaky in both cases.

First, there's the matter of generative AI doing what it's promised to do. Over the past few years we have had escalating prophecies of the technology laying waste to work as we know it. Dario Amodei, the CEO of Anthropic – OpenAI's closest rival – [warned in May last year](#) that AI could wipe out half of all entry-level white-collar jobs within the next five years, for example. "The current generation of AI large language models will not be doing that," Zitron says confidently. "My evidence is they're basically the same as they were a year ago. They have the same efficacy. And every attempt they make to try to turn these into something that can actually do things autonomously has failed." LLMs hallucinate and give wrong answers, they give different answers every time, they cannot really *learn*, or create, or perform a lot of complex tasks, he argues. He questions even describing this technology as "intelligence".

"It's intelligent in the same way a pair of dice are intelligent," he says. "Large language models are transformer-based architecture that use large-scale probability to generate the next token. Now they

do this at scale, so you might think, 'Oh, it's coming up with things.' No, it has a large corpus of data, and so many parameters that it pulls from to generate an output. That is all that is. We would not credit an Excel formula with intelligence, and we should not credit generative AI as intelligent."

Obviously, many people disagree with Zitron, especially when it comes to AI replacing jobs. In industries from film-making to customer service to government agencies to tech itself, insiders say AI tools are enabling them to do the same things with fewer people. Even if it doesn't replace 50% of jobs, its effect on the workplace is likely to be transformative. A survey last June found that entry-level jobs had [dropped by nearly a third](#) in the UK since the launch of ChatGPT.

Zitron argues that "correlation does not equal causation" and points to reports that suggest the role of machine learning in job cuts is either unproven or overstated. A [recent MIT report](#) into the "state of AI in business in 2025", for example, found that 95% of companies attempting to integrate AI in their businesses were getting "zero return". "Most GenAI systems do not retain feedback, adapt to context, or improve over time," it said.

That leads to the second part of Zitron's argument: that the economics of the AI boom just don't stack up. The amounts of money pouring into AI investment are unlike anything the world has ever seen. The "magnificent seven" – Alphabet (parent company of Google), Amazon, Apple, Meta, Microsoft (which owns 27% of OpenAI), Nvidia and Tesla – currently make up [34% of the S&P 500](#), the US stock market index that represents about 50% of the global market. As the dominant manufacturer of GPUs (graphics processing units – the extremely powerful chips on which AI depends), Nvidia is practically "printing money", says Zitron, but at this stage everyone else is borrowing and spending billions they may never recover.

This is the way Silicon Valley startups have always operated, you could say: operate at an initial loss with a view to establishing market share and reaping profits further down the line. But the current disparity between supply and demand is worryingly huge. When it comes to AI, you need to build big and spend big. A typical datacentre requires tens of thousands of GPUs, with each GPU costing upwards of \$50,000 (£37,000). Then you need the software and networking to knit them all together, a giant building on a vast plot of land to put it all in, and huge amounts of electricity and water to run it all. The cost of 1GW of AI datacentre capacity is [estimated at \\$35bn \(£26bn\)](#). As such, the major players in this business are the deep-pocketed "[hyperscalers](#)" like Google, Meta, Amazon, Microsoft and Oracle.

When you look at the demand side, the picture is less rosy, and a lot more hazy. OpenAI alone has committed to [spending \\$1.4tn \(£1tn\)](#) on AI infrastructure over the next five years, for example, but its revenue for 2025 is expected to be about \$20bn (£15.8bn). There seems to be a constant carousel of deals and agreements between AI companies, but when you look at it, says Zitron, much of the time these companies are essentially paying each other. Nvidia, for example, [announced a \\$100bn investment in OpenAI](#) last September; in return, OpenAI will use the cash to buy Nvidia chips. Similar deals abound in this space, as Zitron has forensically detailed. Even with non-magnificent seven "neocloud" companies, like CoreWeave, Lambda and Nebius, which build datacentres then rent out their GPU capacity to others, the bulk of their business is coming from the likes of Google, Microsoft, Amazon and Nvidia, Zitron says. "When you remove the hyperscalers, there's less than a billion dollars total in AI compute revenue in 2025."

As for the profitability, ChatGPT now has an estimated 800 million users, but the vast majority of them are paying nothing. Even for paying subscribers, "when you connect a user to an AI model like GPT, each thing the user does varies in expense vastly. A user could ask a very simple question, or they could ask a question that the model interprets as needing a complex answer," Zitron says. There are no economies of scale here; each question requires "compute" – as in computer processing activity – at the supplier's expense. "The more someone is a power user of these platforms, the more they're going to cost you. This is almost the inverse of how the valley works." And if the answer is not satisfactory and must be reformulated, as is often the case, "that's more compute burned, making you no extra money". AI models are getting cheaper and more sophisticated all the time, we are told, but only by using more compute. "It's like the price of petrol going down a bit, but you have to drive another 250 miles to get somewhere. So this is really problematic – because it means that there is no profitability point."

Again, none of this means the great AI crash will happen, but "if I'm wrong, I don't know how I'm wrong," he says. "Every counter I've read to my work is mostly just wishcasting of 'and then the AI gets better'."

Many have accused Zitron of having an axe to grind against big tech, but he refutes that: "I have an axe to grind against those who don't want to talk about reality." He certainly doesn't shy away from attention, but that's not why he got into this business, he explains. "I like writing. I like pulling things apart. I like solving puzzles. I guess I like being able to understand things. A lot of this is just me trying to explain it to myself, rather than an audience." He had no formal training in economics or computer science and has never worked in tech. "I've learned basically everything from the ground up."

Zitron has, it seems, always been technological though. He has built 10 personal computers over his lifetime, he says. It started when his father bought him a PC card with a dial-up connection when he was 10. "So I was online from quite an early age. I immediately was just like, 'This is the future. I adore this. I love that I can talk to people and game with people.' I was quite a solitary child. I didn't have a lot of friends, but I made a lot of friends online."

Growing up in Hammersmith, west London, his parents were loving and supportive, Zitron says. His father was a management consultant; his mother raised him and his three elder siblings. But "secondary school was very bad for me, and that's about as much as I'll go into." He has dyspraxia – a coordination disability – and he was diagnosed with ADHD in his 20s. "I think I failed every language and every science, and I didn't do brilliant at maths," he says. "But I've always been an asshole over the details."

After studying media and communications at Aberystwyth University, he began writing for gaming magazines, but "I got to a point where I was miserable in London." So he relocated to New York in 2008 and began working in tech PR. He can't contemplate returning to the UK, he says. He doesn't talk about his personal life beyond saying he has a son, which is why he lives in Las Vegas. He doesn't mind it there: "Everyone's weird so no one's weird." It has been reported that he is twice married and twice divorced.

Zitron continues to work in tech PR, which seems jarringly at odds with his career as a tech agitator – either like biting the hand that feeds him or a conflict of interest. He doesn't see it like that. He doesn't have AI clients, or work with big tech, he says, and only has a few clients now. The work has given him a network of contacts in the industry, and possibly helped him to market himself (in 2013 he published a book titled *This Is How You Pitch: How To Kick Ass in Your First Years of PR*). He may not be doing the PR stuff much longer, though. The media side of things is "making up more of my income these days than I ever expected it to". He's writing a new book, due out next year, called *Why Everything Stopped Working*. "It's kind of a dig into how the world got the way it did and technology is everything now." Just one chapter is about AI, he adds.

If Zitron does have an axe to grind, it's against neoliberal capitalism in general: "I don't think people have taken seriously enough how bad deregulation of financial markets, by Thatcher, by Reagan, was. I don't think people take seriously enough how bad it was not putting people in prison for the great financial crisis ... I don't think people have taken seriously the threat of growth-focused capitalism and growth at all cost."

Rather than leading us to a utopian future, Zitron sees AI as the logical conclusion of neoliberalism. "The biggest thing we've learned from the large language model generation is how many people are

excited to replace human beings, and how many people just don't understand labour of any kind," he says.

Zitron is no longer quite so alone in his assessment. He's on the same page as Cory Doctorow, for example, who has appeared on his podcast, and whose "enshittification" thesis also alleges that tech companies are now more motivated by profit than making more useful products. Meanwhile, other AI sceptics, such as cognitive scientist Gary Marcus, [complain](#) they have been making the same arguments as Zitron "but in his narrative, I don't exist". Either way, the backlash to AI is building: local groups are opposing the construction of environmentally destructive datacentres; consumers are grating against the insertion of AI into every conceivable product; creators are taking legal action against the industry's theft of their work; there is public outrage over social media harms, epitomised by Elon Musk's Grok creating [nonconsensual borderline-deepfake porn](#).

Meanwhile, speculation about the AI bubble bursting continues to grow. Now everyone from the [Bank of England](#) to [Microsoft boss Satya Nadella](#) are raising the alarm. Michael "Big Short" Burry says he is [betting against Nvidia](#), and recently the New York Times ran an op-ed speculating that [OpenAI will run out of money](#) within the next 18 months. It could be sooner than that, Zitron thinks: this month, the big tech companies start reporting their annual earnings for 2025. Most of them have been cagey about their revenues from AI specifically, he says. "Why would they do that? Well, because they're not very big. So this whole thing is – to use a phrase I hate – it is a *vibe*." If something serious happens, like Nvidia missing its targets, it could prompt a rethink of the whole sector, and possibly a new global financial crisis. All those datacentres might well end up as empty shells. Ultimately, we could be witnessing "the largest laser-tag arena construction of all time," he jokes.

Zitron doesn't actually enjoy being contrarian, he insists. "It isn't fun being alone in an idea, which is actually why I think a lot of people are pro-AI, because it's much easier to do that."

He doesn't hate tech, or even AI, he says. "I love technology, but I hate what the tech industry is doing ... If you can't critique this stuff without it being claimed that you don't support the world or innovation, I think you realise we're in this weird peasant economy where even wealthy, well-to-do famous people have to kneel at the feet of these companies. And these companies have done very little to make our lives better, all while making so much more money than we will ever have."

He just wants to tell it like it is. "It'd be much easier to just write mythology and fan fiction about what AI could do. What I want to do is understand the truth."

Produced by wallabag with tcpdf

Please open [an issue](#) if you have trouble with the display of this E-Book on your device.