
Silicon Valley Is Making Its Wildest Promises Ever. This Number From Anthropic Proves It.

Published by

Estimated reading time

6 min

Added on

2026-09-02

Address

<https://slate.com/technology/2026/09/artificial-intelligence-anthropic-spacex-ipo.html>

Silicon Valley Is Making Its Wildest Promises Ever. This Number From Anthropic Proves It.

[Sign up for the Slatest](#) to get the most insightful analysis, criticism, and advice out there, delivered to your inbox daily.

When Elon Musk was getting ready to take SpaceX public in June, the company advertised, in its government paperwork, the largest “total addressable market” in human history: \$28.5 trillion per year, with \$26.5 trillion (lol) attributed to artificial intelligence. A TAM is a company’s assessment of all the money it could theoretically capture in every market it could serve. For an ice cream truck visiting a summer camp, the TAM is all the kids at the camp multiplied by how much ice cream each kid would buy that summer. A big TAM tells investors that capturing even a sliver of the market could mean a great business. SpaceX’s claim is that its total addressable market is [nearly 90 percent of the United States’ gross domestic product](#).

This \$28.5 trillion claim is a *lot*. Of course, Musk is a serial fabulist [who says tons of things that don’t pan out](#), and he’s also [functionally not subject to the country’s laws](#) while acting as the donor kingpin of Republican politics during a Republican presidency. But this [\\$26.5 trillion claim](#) was still interesting, because it’s not like that immunity from the law applies to all the lawyers and bankers doing the paperwork for him. *They* can’t safely lie on Securities and Exchange Commission forms, and yet here those forms were.

That realization hit again last week, when the Wall Street Journal reported that A.I. rival Anthropic will [project a total addressable market of \\$30 trillion](#). (The company already confidentially [submitted its forms](#), so we’ll probably know soon.) SpaceX had barely finished making the most exorbitant market claim in capitalist history before word leaked that Anthropic would make a bigger one. “A.I. company makes big, brash economic claim” isn’t surprising, but these stories were striking because of the medium involved. You can maybe get away with a lot of bullshitting in panel discussions and press releases, but these are *documents filed with the SEC*, and the titans of A.I. are not chilling out in this venue either. And whatever Musk can get away with because federal authorities now do whatever Donald Trump wants, Anthropic’s situation is the reverse. This company [gets special treatment the other way](#) because the big man [doesn’t like them](#).

In theory, going public should be the moment when A.I. companies start talking in slightly more restrained terms so that they aren’t sitting ducks for prosecution and civil suits. But rather than dialing anything back, the big players are just revving their engines more and more. This seems to present two options: You really can just say anything in an SEC filing now, or the barons of A.I. are somewhere between high on their own supply and under the sincere impression that they’re about to own the world.

There’s room for marketing and hype here, sure. The SEC reviews paperwork to make sure it follows disclosure rules but doesn’t weigh in on its accuracy. An estimate of a market size is not a promise to take over 100 percent of the market, and that leaves room for an over-the-top pitch. But investors can sue over “material” falsehoods and omissions, and they often do. Calling something an estimate doesn’t shield a company from scrutiny either. Not even two years ago, a much different SEC charged and settled with the biotech company Zymogen after the company claimed a \$1 billion TAM while, the SEC said, its internal teams were projecting a market many times smaller than that. Former executives and insiders agreed to a \$125 million settlement in [a related class](#)

[action](#). All of this for a company that was way off on a \$1 billion TAM projection. It makes you wonder what might support *tens of trillions*.

In SpaceX's case, Musk's people at least laid out their \$28.5 trillion "math," and you could chuckle at how silly it was. Most of it was a \$22.7 trillion estimate for A.I. enterprise applications, based on the value of the entire digital economy. (SpaceX admits that its A.I. cannot yet address this market in full.) There was also a TAM of \$600 billion per year in "digital advertising," which the company noted was the entire world's digital advertising market in 2025 outside of Russia and China. SpaceX counted all of that as a market it could compete for, because the company owns what used to be Twitter. OK! That's Elon Musk!

Anthropic will have to get to its own market projection without owning a declining microblogging platform, and until its documents are public, we won't know how the company justifies it. The Journal has offered a preview, saying the company is "looking at the full scope of work that could be completed with AI models." Anthropic will also have to figure out how to sell A.I. without burning cash. (Anthropic's description of its current profitability will be *very* interesting.) But these guys are capturing a ridiculous volume of cash: [an annualized rate of \\$47 billion as of May](#), which isn't the same as money in the bank but could also undersell things if the company has grown. I think that if you are still insisting that every A.I. business will go up on smoke, it's time for an update. You are wishcasting rather than looking at the world as it is.

Hardened priors go both ways, though. Anthropic is a real business that needs to justify extraordinary and frankly outrageous valuations and infrastructure spending to keep the party going. Projecting a market of (approximately) infinity dollars is one way to do that. Reuters [noted that the big projection](#) would "anchor the growth story behind its planned IPO," a story the rest of us aren't obligated to tell.

But it *is* a new development that these companies are going public and now telling those growth stories in ways that are, in theory, more legally reviewable. A.I.'s most famous people have habits of making big, bold assertions about the technology and then softening them over time. However, these walkbacks have tended to be easy to make, because they've been in response to A.I. being less disruptive than they first projected. Sam Altman acknowledged on a podcast this week that "we've all been too ambitious on the timelines," but he was able [to talk softly about how that was a good thing](#). In the spring, [he explained](#) he was "delighted to be wrong" about A.I. taking more people's jobs. The past half-decade of lofty A.I. statements has come without many consequences, because when they've been wrong, they've been wrong in ways that seem good for society. They've often come from people at private companies whose investors were not going to sue them even if they thought they had a good reason.

This is a slightly different frontier, where big claims are tied to newly public companies that could cause immediate havoc if they do poorly. Anthropic's forthcoming projection seems to be that it can be unlike anything that's ever existed in business. If a company tells the world its total addressable market is \$30 trillion, it might not have to bring in *\$30 trillion* to look like an honest broker. But it should probably become one of the most impressive businesses ever, right? If not, wouldn't that invite pretty harsh recriminations from either aggrieved investors or motivated regulators and prosecutors? At the least, those responses seem more likely than if the boss were tossing around vagaries on a podcast.

Produced by wallabag with tcpdf

Please open [an issue](#) if you have trouble with the display of this E-Book on your device.